Colorado River District/Courtesy photo

NextGen Research DataStream: Community Contributions Towards Improved Hydrologic Predictions

Lynker: Jordan J. Laser, Zach Wills, Nels Frazier
Alabama Water Institute: James Halgren, Arpita Patel

Road Map

1. DataStreamCLI Motivation
2. DataStreamCLI Design
3. Research DataStream Motivation
4. Research DataStream Design
5. Research DataStream State
6. Hands-on workshop

Colorado River District/Courtesy photo

DataStreamCLI Motivation

Workflow that abstracts the laborious process of collecting input data and executing NextGen

→ DataStreamCLI will perform every step in executing NextGen. The user can compute steps separately and provide those files directly. *Batteries included but flexible*

Reproducibility

→ Enforces the NextGen In A Box standard run directory.

Metadata – all relevant information about user inputs, code versions, host architecture, etc. so that a NextGen execution can be understood and reproduced.

Need baseline NextGen dataset to evaluate new realizations against

Scaling from laptop to HPC

→ DataStreamCLI can scale from a laptop to the cloud. DataStreamCLI performs the on-server workflow that drives the Research DataStream.

1 Prepare configuration and data inputs

2 Validation

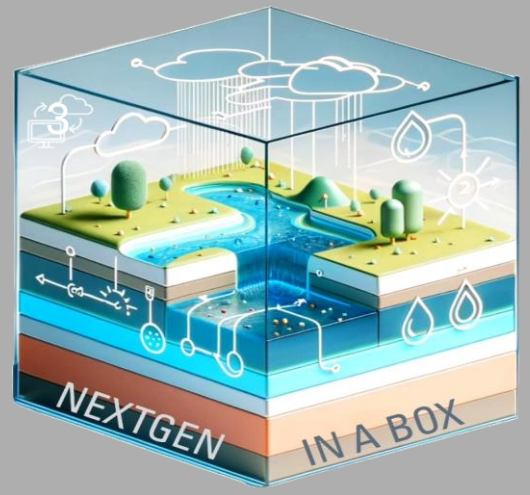

4 Merkle Tree

DataStreamCLI

DataStreamCLI refers to the software chain that builds and validates NextGen input packages (ngen-run/), executes NextGen through NextGen In A Box (NGIAB), and versions the entire run for reproducibility.

This enforces a standard folder (ngen-run/), which enables interoperability and reproducibility.

ngen-run/

#	name	type	size
0	config	dir	288 B
1	forcings	dir	343.8 KiB
2	lakeout	dir	64 B
3	outputs	dir	592.3 KiB
4	restart	dir	64 B

Lynker NOAA CIROH ALABAMA Alabama Water Institute

Required steps to build ngen-run/config and ngen-run/forcings

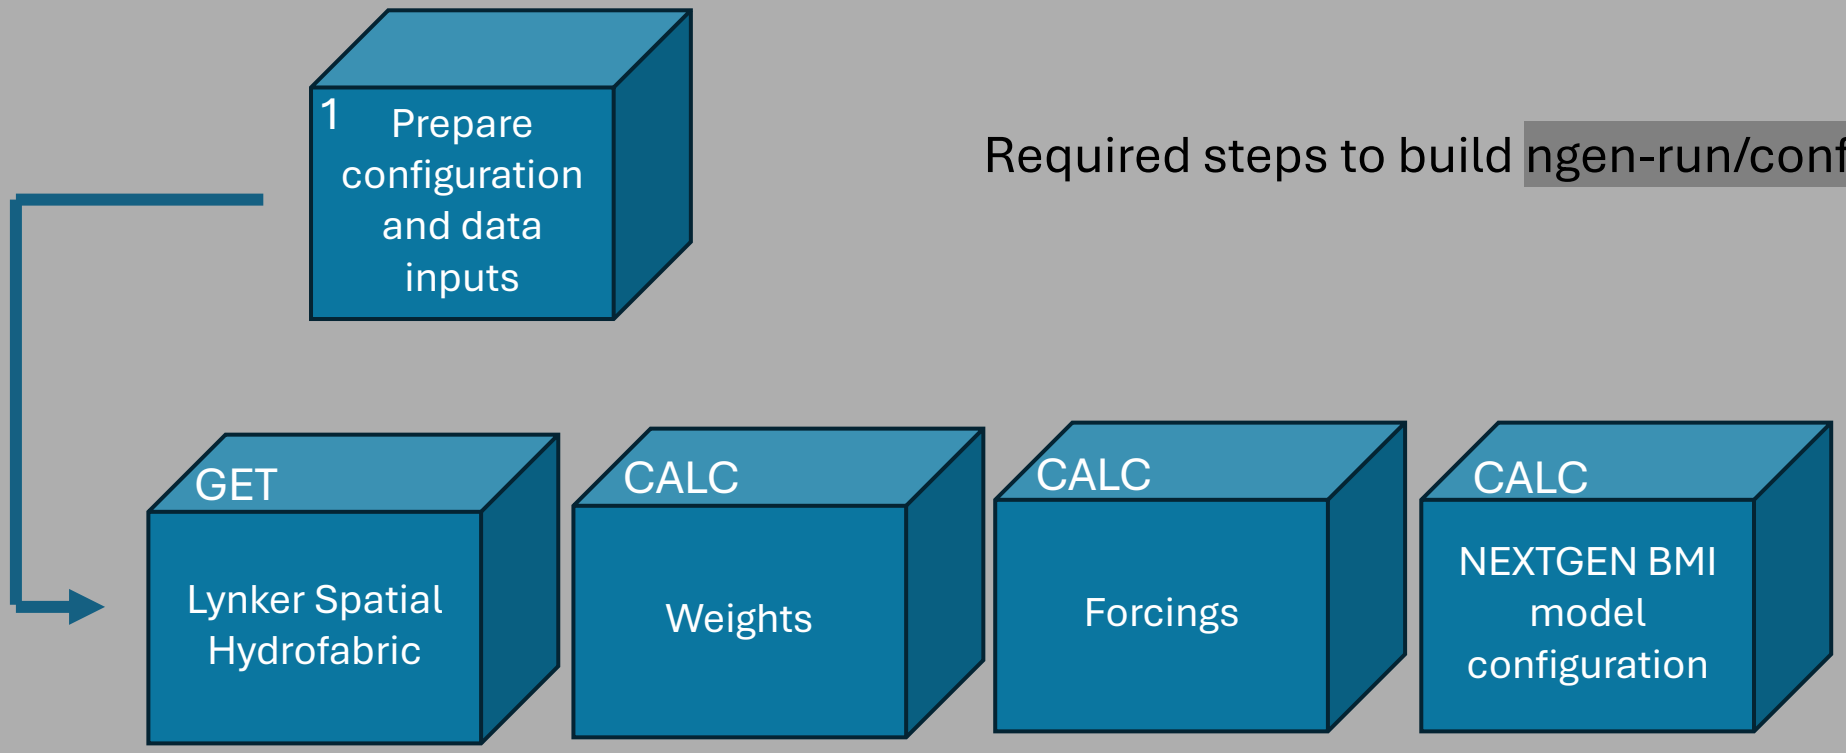

Defines spatial domain

Indices and coverage used to extract catchment averaged forcings. Calculated by exactextract.

Performs conversion between National Water Model and NextGen forcings formats

Required files for NextGen BMI modules

ngen-run/

#	name	type	size
0	config	dir	288 B
1	forcings	dir	343.8 KiB
2	lakeout	dir	64 B
3	outputs	dir	592.3 KiB
4	restart	dir	64 B

LynkerSpatial

Open Data, Open Science

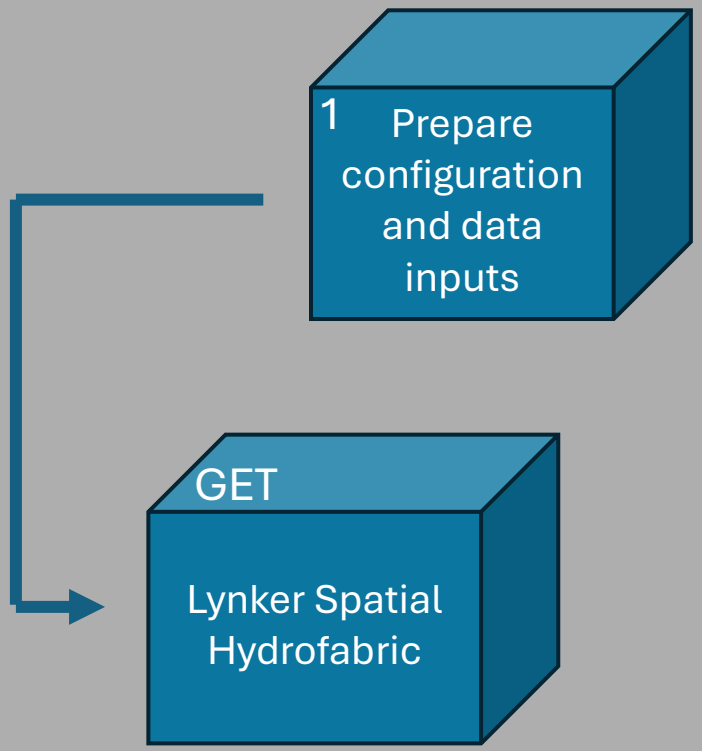

Defines spatial domain

ngen-run/

#	name	type	size
0	config	dir	288 B
1	forcings	dir	343.8 KiB
2	lakeout	dir	64 B
3	outputs	dir	592.3 KiB
4	restart	dir	64 B

Lynker NOAA CIROH ALABAMA Alabama Water Institute

1 Prepare configuration and data inputs

GET
Hydrofabric

CALC
Weights

Defines spatial domain

Indices and coverage used to extract catchment averaged forcings

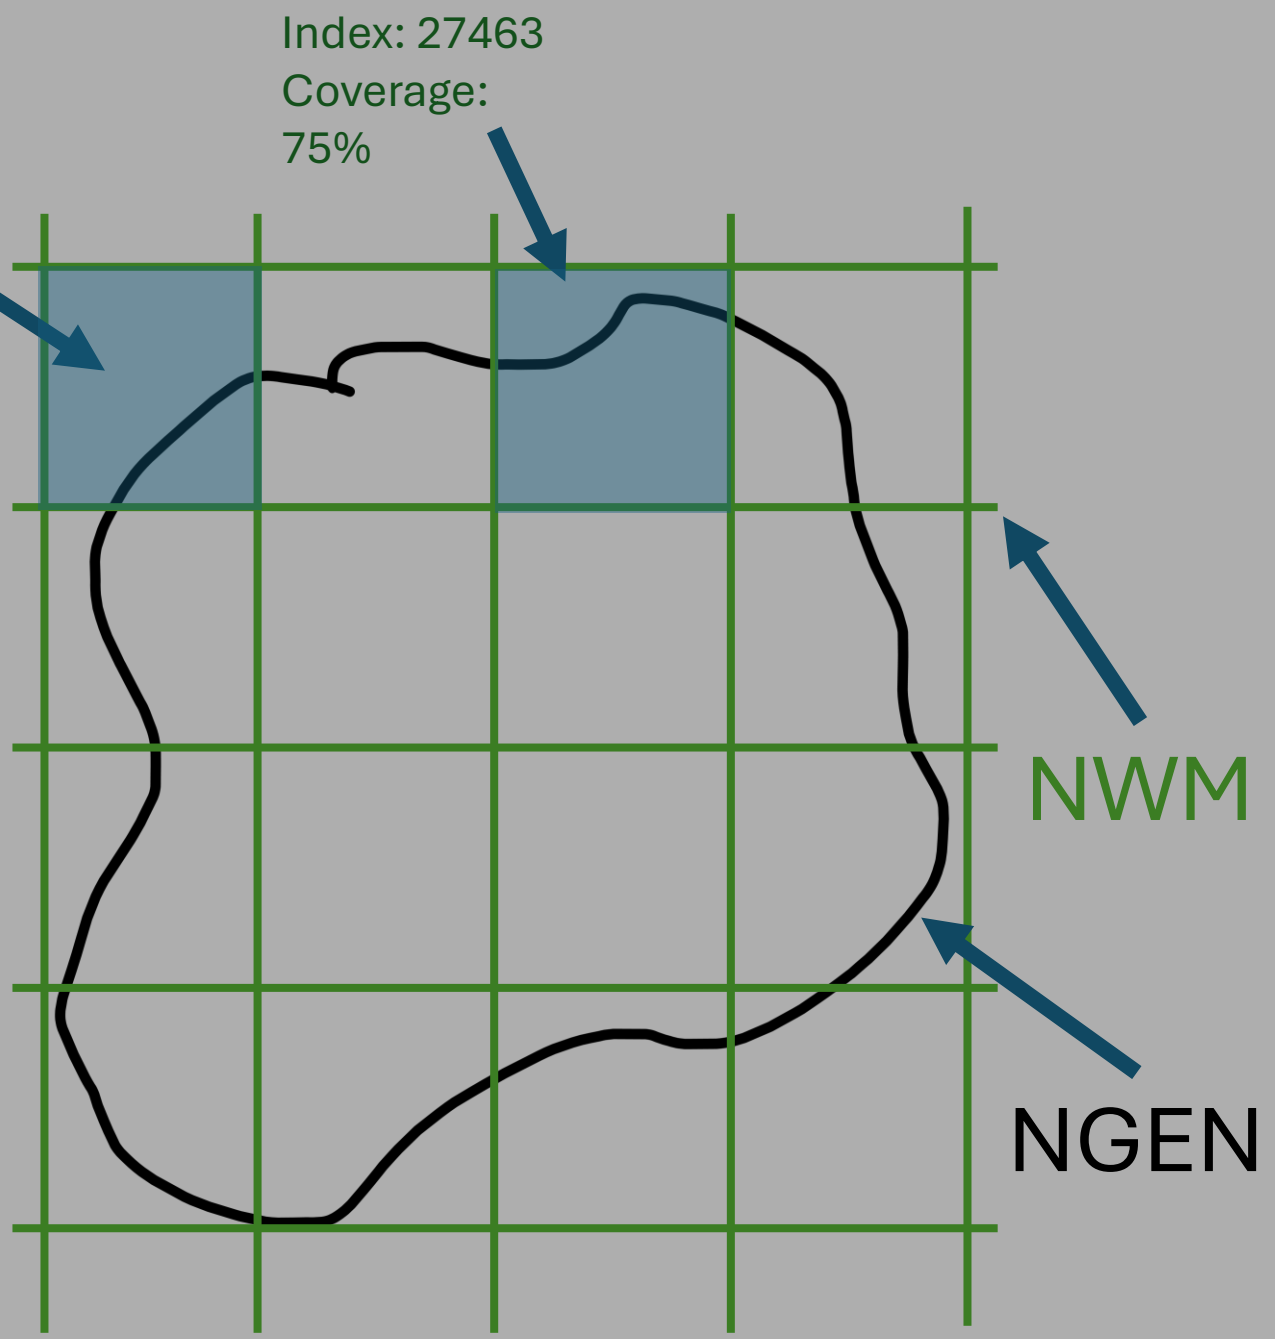

1 Prepare configuration and data inputs

GET
Hydrofabric

Defines spatial domain

CALC
Weights

Indices and coverage used to extract catchment averaged forcings

CALC
Forcings

Performs conversion between National Water Model and NextGen forcings formats

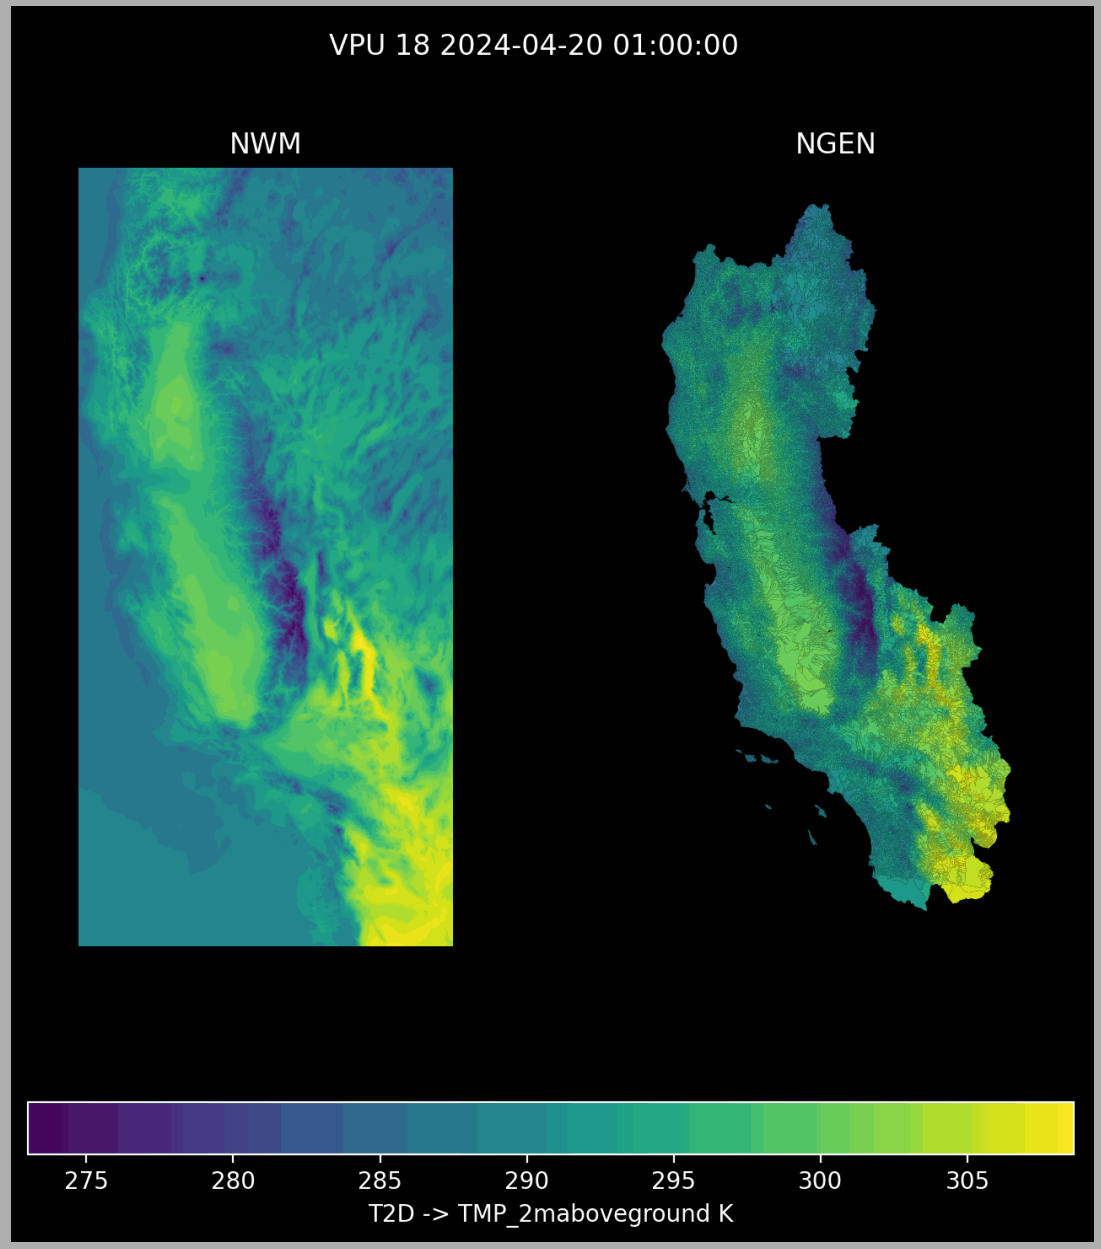

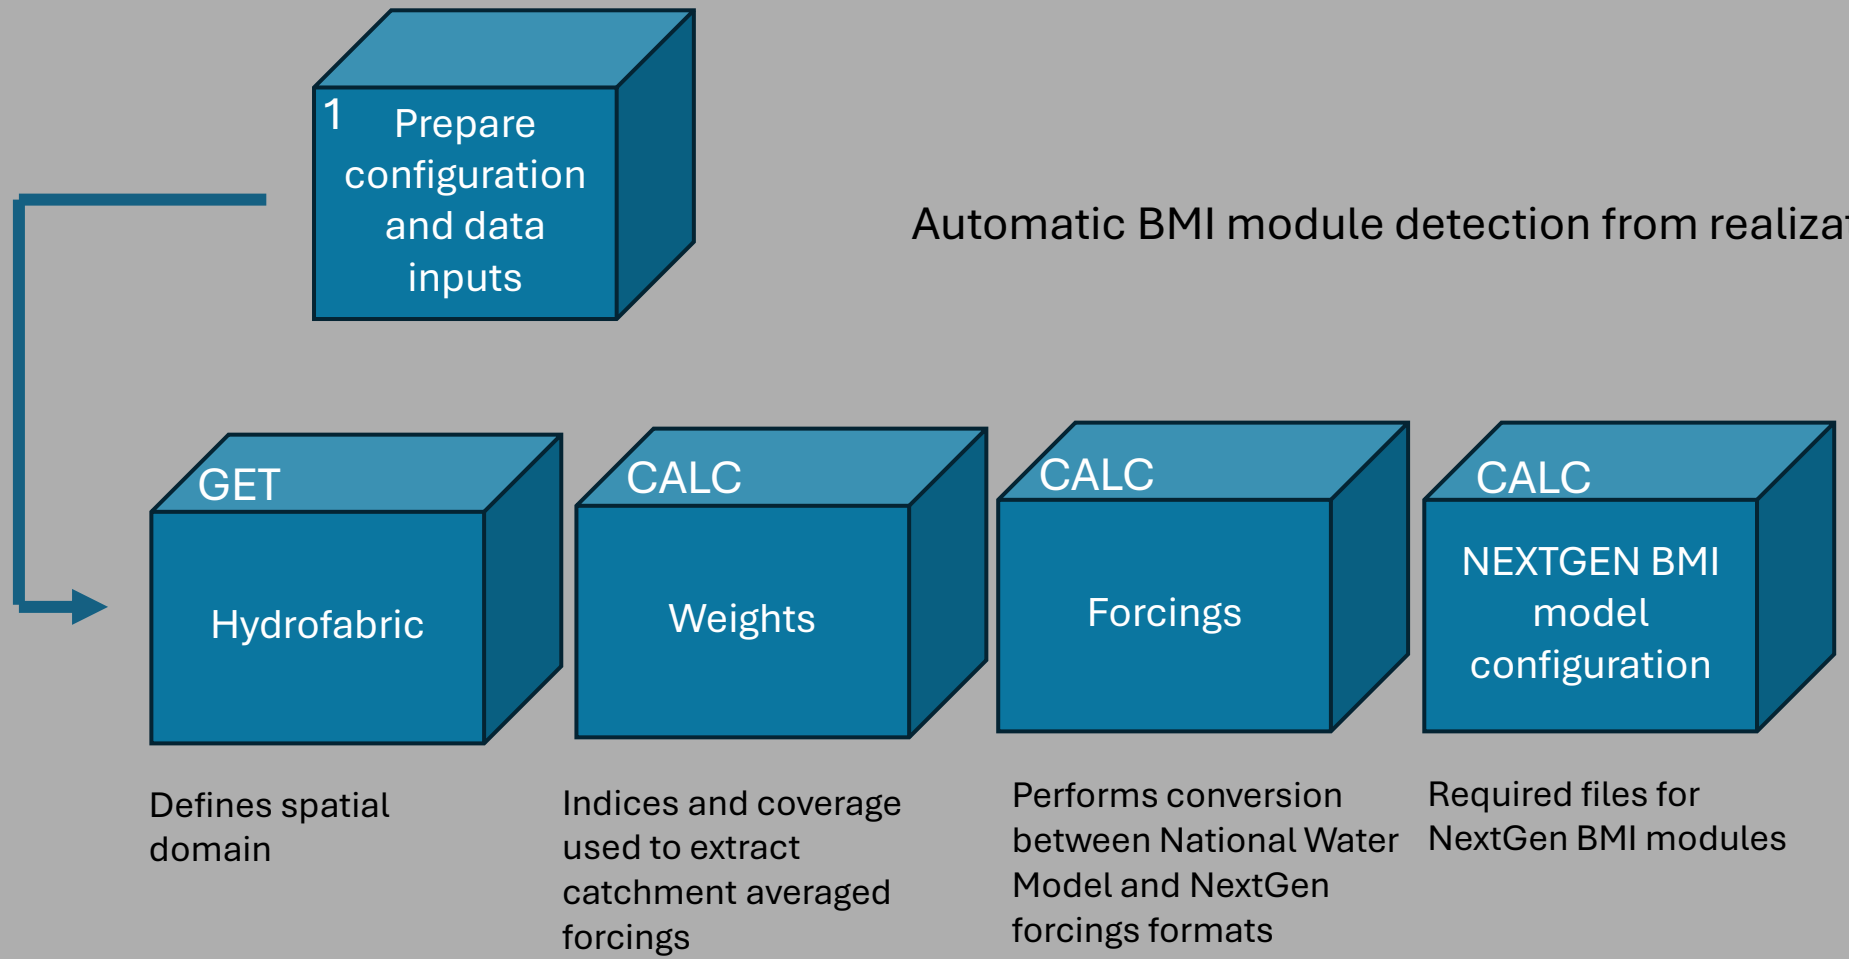

Automatic BMI module detection from realization file

Defines spatial domain

Indices and coverage used to extract catchment averaged forcings

Performs conversion between National Water Model and NextGen forcings formats

Required files for NextGen BMI modules

- Supported BMI module config generation
  - PET, CFE, Noah-OWP-Modular, t-route
- Coming soon
  - SoilFreezeThaw, TopModel, LSTM, others

The resource directory is used as a cache for files that can be reused.

For a given domain, hydrofabric, weights, and BMI config files can be reused.

For a given domain and time, forcings can also be reused.

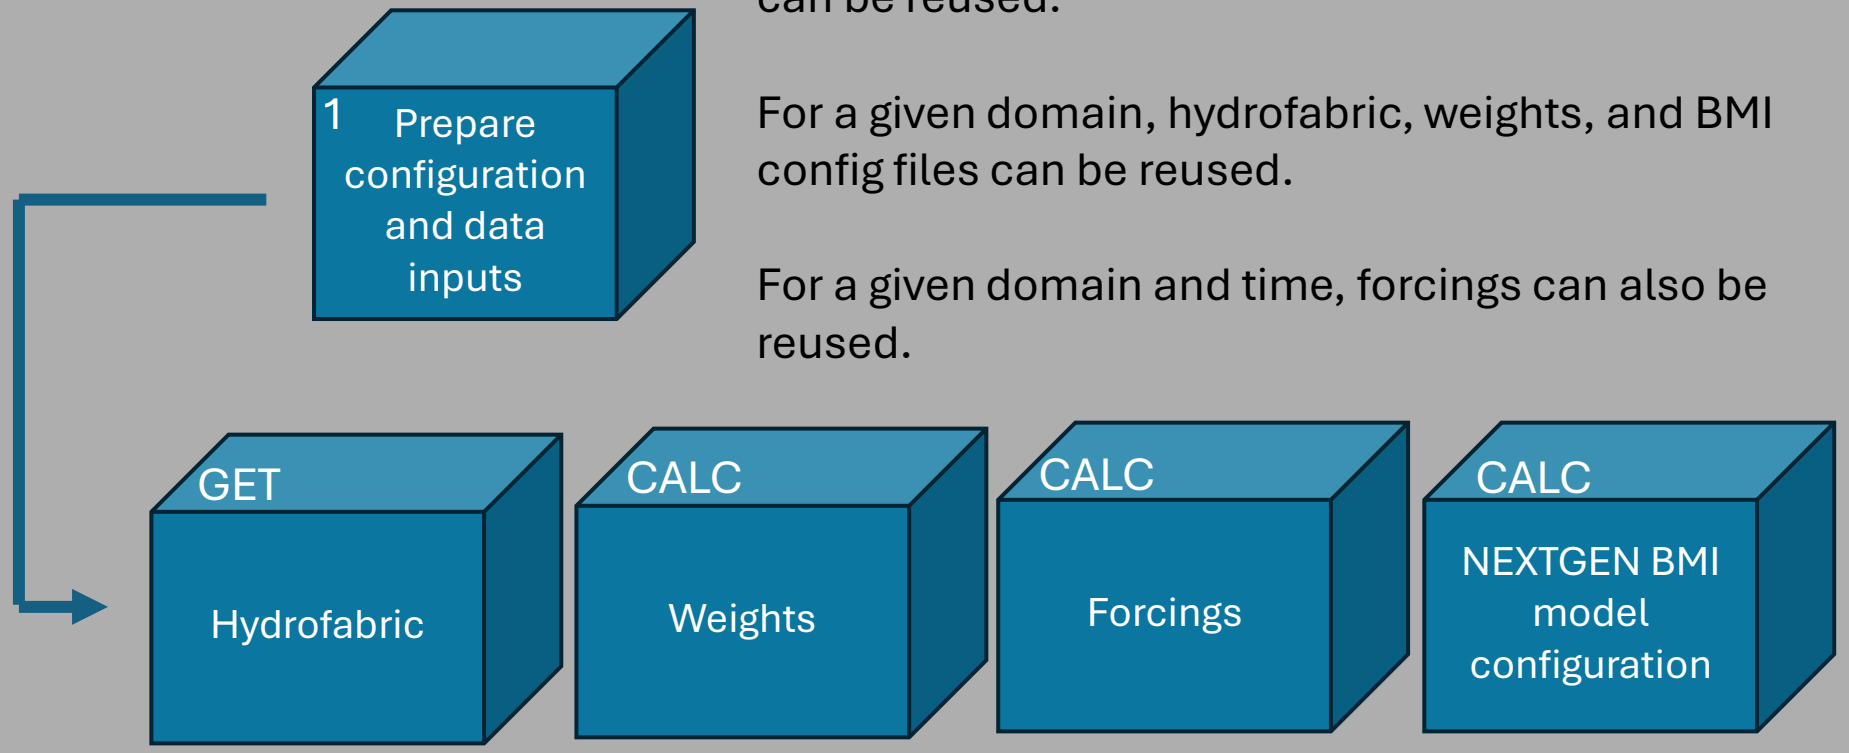

Defines spatial domain

Indices and coverage used to extract catchment averaged forcings

Performs conversion between National Water Model and NextGen forcings formats

Required files for NextGen BMI modules

```
RESOURCE_DIR/
|
├── config/
|   |
|   ├── ngen-bmi-configs.tar.gz
|   |
|   ├── realization.json
|
├── datastream
|   |
|   ├── partitions.json
|   |
|   ├── weights.json
|
|── hydrofabric
|   |
|   |── nextgen_01.gpkg
|   |
|   ├── nextgen_01.parquet
|   |
|   ├── weights.parquet
|
├── nwm-forcings/
|   |
|   ├── nwm.t00z.medium_range.forcing.f001.conus
|   |
|   ├── ...
|
├── ngen-forcings/
|   |
|   ├── forcings.tar.gz
|
```

https://github.com/CIROH-UA/ngen-datastream/tree/main?tab=readme-ov-file#resource_dir-datastream-resour

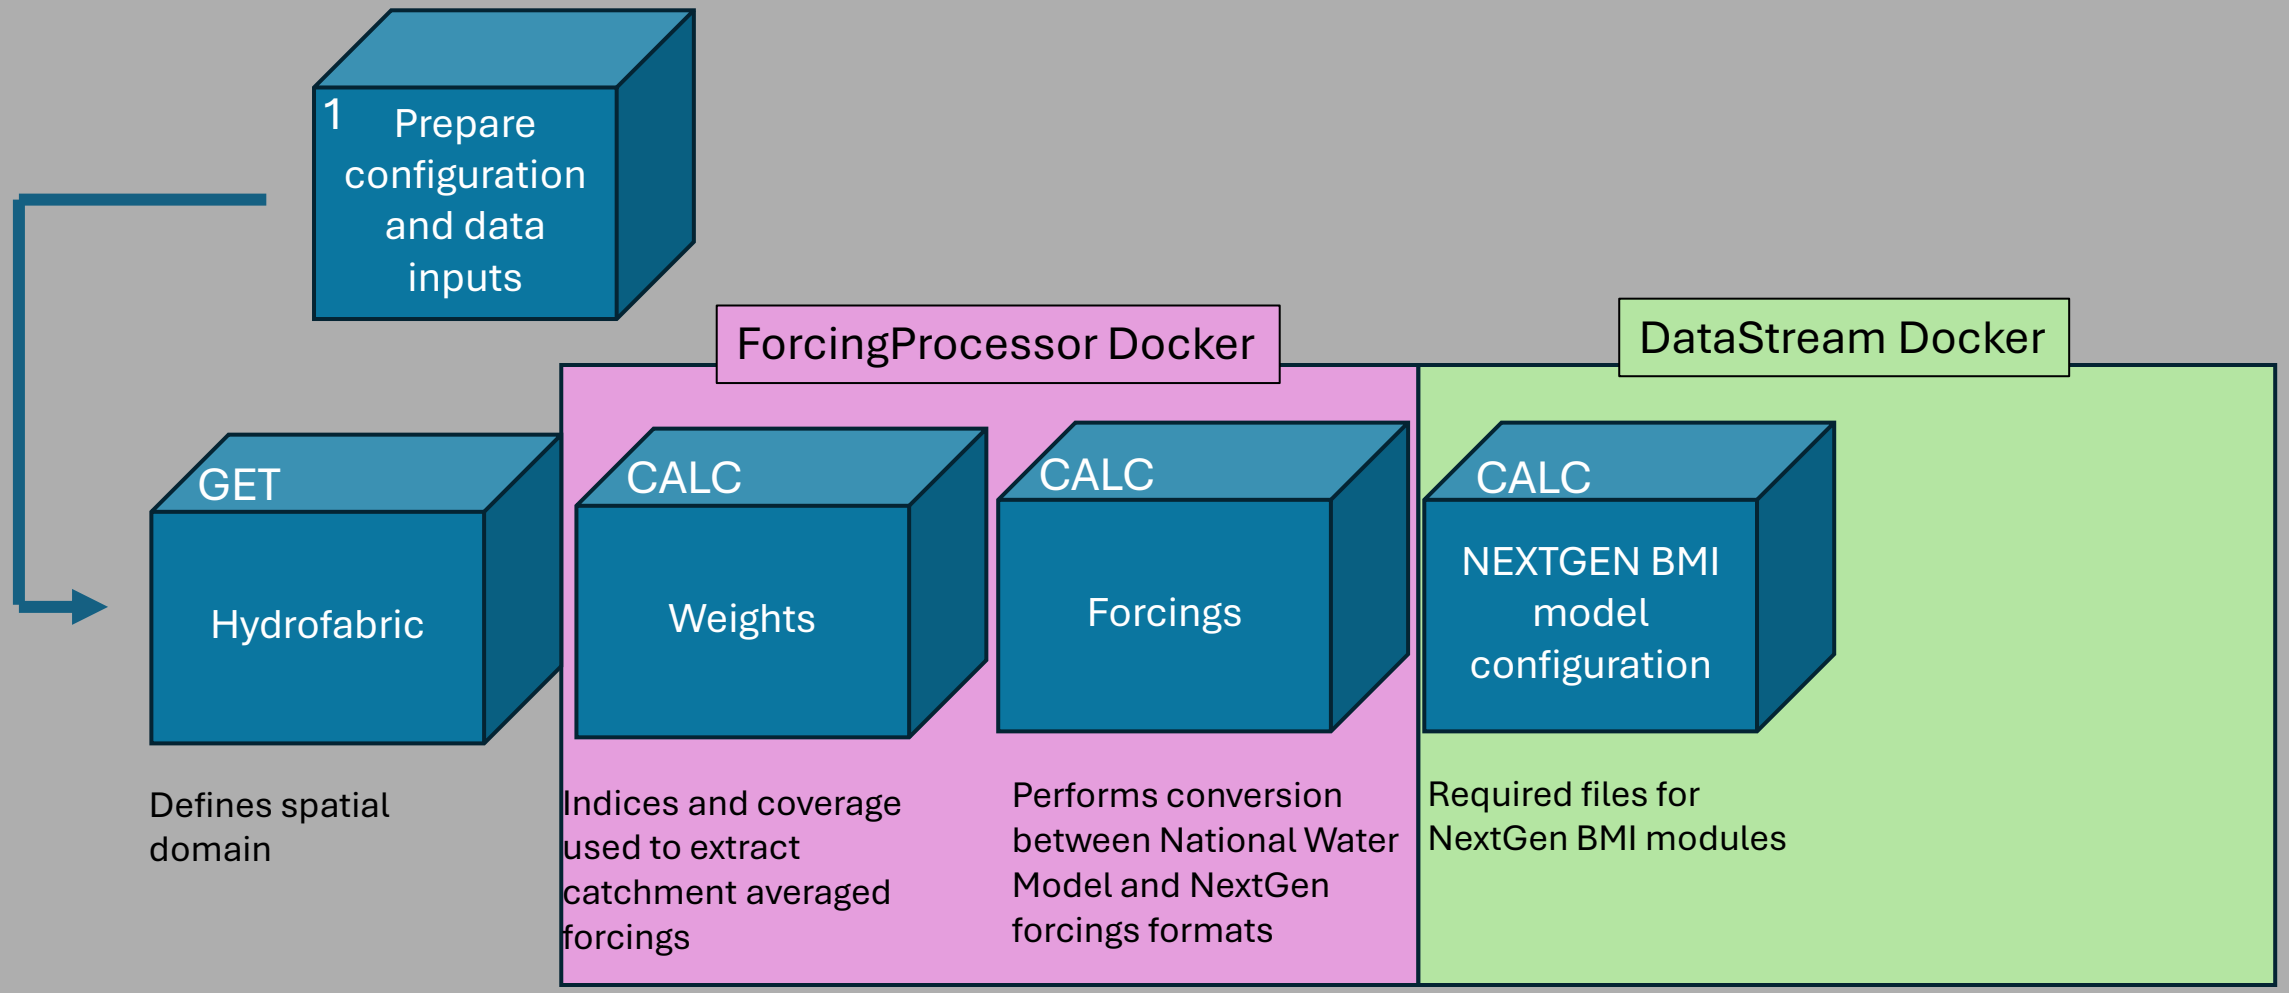
1 Prepare configuration and data inputs
ForcingProcessor Docker
DataStream Docker
GET
Hydrofabric
CALC
Weights
CALC
Forcings
CALC
NEXTGEN BMI model configuration
Defines spatial domain
Indices and coverage used to extract catchment averaged forcings
Performs conversion between National Water Model and NextGen forcings formats
Required files for NextGen BMI modules

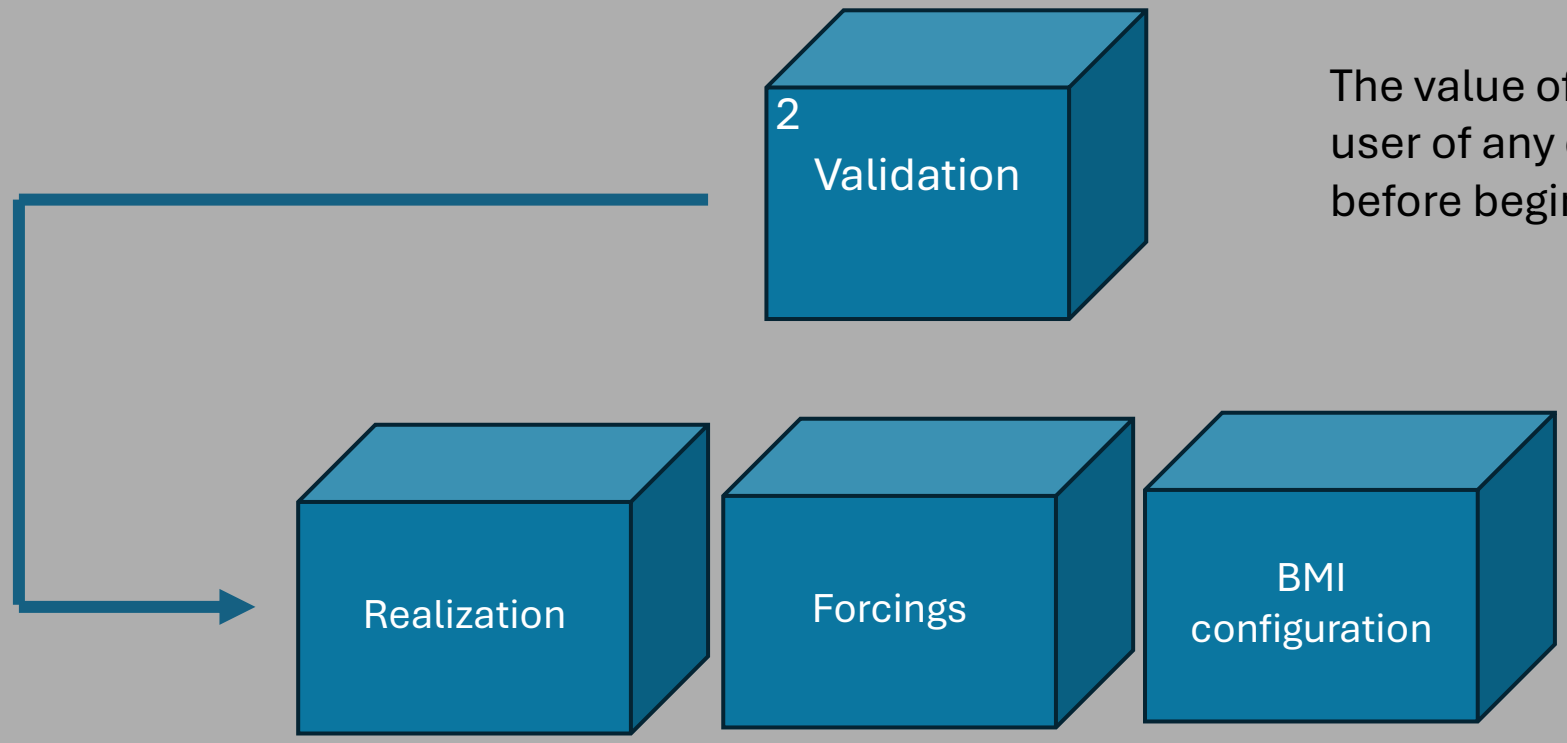

The value of the validation step is to notify the user of any errors in the NEXTGEN run package before beginning the execution

Coming soon -> BMI module variable mapping validation

Required for

Ensures the user has supplied a valid realization file to configure NEXTGEN

Ensures a forcing file exists for each catchment in the hydrofabric and for each time step specified in the realization

Ensures all BMI model configuration files exist.

ngen-run/

#	name	type	size
0	config	dir	288 B
1	forcings	dir	343.8 KiB
2	lakeout	dir	64 B
3	outputs	dir	592.3 KiB
4	restart	dir	64 B

Lynker NOAA CIROH ALABAMA Alabama Water Institute

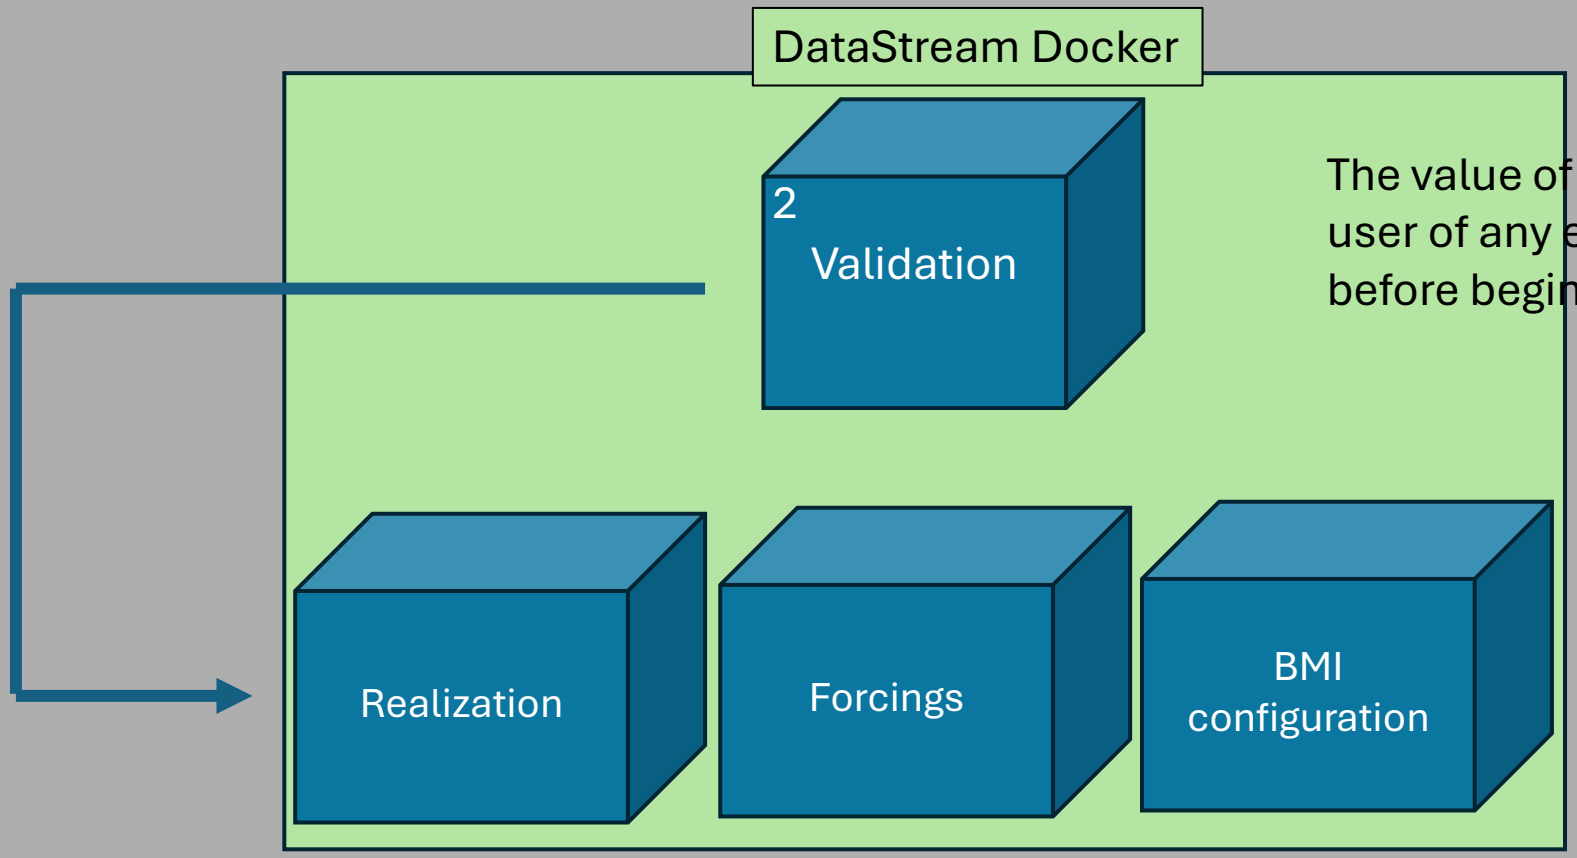

The value of the validation step is to notify the user of any errors in the NEXTGEN run package before beginning the execution

Coming soon -> BMI module variable mapping validation

Required for

Ensures the user has supplied a valid realization file to configure NEXTGEN

Ensures a forcing file exists for each catchment in the hydrofabric and for each time step specified in the realization

Ensures all BMI model configuration files exist.

ngen-run/

#	name	type	size
0	config	dir	288 B
1	forcings	dir	343.8 KiB
2	lakeout	dir	64 B
3	outputs	dir	592.3 KiB
4	restart	dir	64 B

Lynker

1 Prepare configuration and data inputs

2 Validation

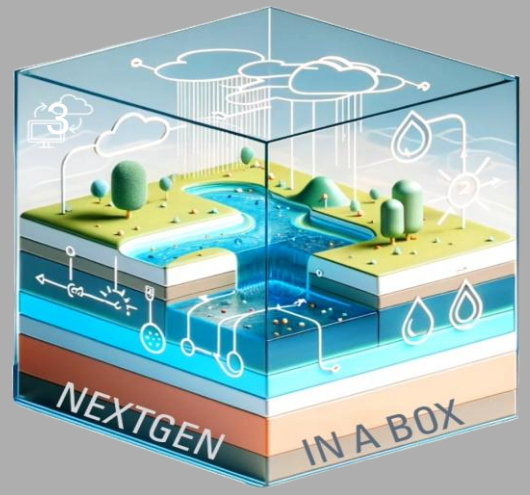

4 Merkle Tree

DataStreamCLI

DataStreamCLI refers to the software chain that builds and validates NextGen input packages (ngen-run/), executes NextGen through NextGen In A Box (NGIAB), and versions the entire run for reproducibility.

This enforces a standard folder (ngen-run/), which enables interoperability and reproducibility.

ngen-run/

#	name	type	size
0	config	dir	288 B
1	forcings	dir	343.8 KiB
2	lakeout	dir	64 B
3	outputs	dir	592.3 KiB
4	restart	dir	64 B

Lynker NOAA CIROH ALABAMA Alabama Water Institute

- Merkle Tree based hashing algorithm
  - “Root” hash allows for quickly identifying if two ngen-run directories are different.
  - Ability to query whether some file is a part of the tree represented by the root hash
  - Ability to compare files without opening them


```
[jlaser@LYNK-59WW6S3 ngen-datastream]$ docker run --rm -v $(pwd)/data/datast
ream_test_VPU09_0520_with_resources_new_realization:/mounted_dir zwills/merk
dir /merkdir/merkdir verify-file -t /mounted_dir/merkdir.file -n "ngen-run/c
onfig/realization.json"
OK: file is still verified by this Merkle tree
```

ngen-run/

#	name	type	size
0	config	dir	288 B
1	forcings	dir	343.8 KiB
2	lakeout	dir	64 B
3	outputs	dir	592.3 KiB
4	restart	dir	64 B

1 Prepare configuration and data inputs

2 Validation

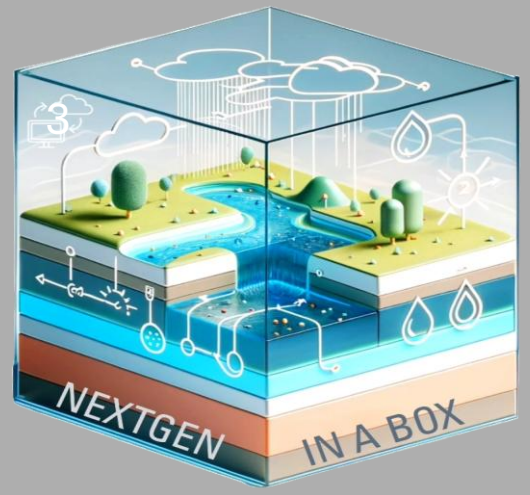

4 Merkle Tree

DataStreamCLI

DataStreamCLI refers to the software chain that builds and validates NextGen input packages (ngen-run/), executes NextGen through NextGen In A Box (NGIAB), and versions the entire run for reproducibility.

This enforces a standard folder (ngen-run/), which enables interoperability and reproducibility.

ngen-run/

#	name	type	size
0	config	dir	288 B
1	forcings	dir	343.8 KiB
2	lakeout	dir	64 B
3	outputs	dir	592.3 KiB
4	restart	dir	64 B

Lynker NOAA CIROH ALABAMA Alabama Water Institute

3. Research DataStream Motivation

Need for regionalized parameterization and modeling

→ Within the NextGen Framework, models and parameters can be configured individually for each catchment

High on-premise cost of HPC

→ Cloud providers offer a cost-effective alternative to purchasing and maintaining expensive on-premise hardware.

The Research DataStream is written in Terraform and made publicly available.

Research 2 Operations

→ The Research DataStream is open to community contributions.

An evaluation workflow is under development to ensure continued improvement of the system.

Lynker

Colorado River District/Courtesy photo

Colorado River District/Courtesy photo

4. Research DataStream Design

- CONUS wide
  - Distributed processing by Vector Processing Unit (VPU)
  - Regional hydrologic processes map to compute resources
  - Outputs are delineated by VPU
- Mimic NWM forecast cycles
  - Short range (18 hourly time steps) 24 times per day
  - Medium range (240 hourly time steps) 4 times per day
  - Analysis assim extend (28 hourly time steps) 1 per day
- Publicly available and editable NextGen configuration files
  - Automated evaluation drives improvement.
- AWS Step Functions state machine
  - Manages infrastructure workflow
- DataStreamCLI
  - Manages on-server workflow

Colorado River District/Courtesy photo

4. Research DataStream State

- VPUs available : 02, 03N, 03S, 03W, 04, 05, 06, 08, 09, 10L, 10U, 11, 12, 13, 14, 15, 16, 18
  - (05, 10L, 10U, 11 not available for medium range)
- Run Types –
  - short range (all initialization cycles),
  - medium range (all cycles, 1st member),
  - analysis assim extend
- Cold start
- NextGen configuration - NOAH-OWP, PET, CFE, and troute.
  - Dynamically read on each execution from publicly available realizations that now hold mutable community parameters.

Colorado River District/Courtesy photo

Research DataStream: Workshop

- https://github.com/CIROH-UA/ngen-datastream/blob/main/docs/CIROH_devcon_2025/workshop.md

- Ask at least 1 anonymous question.

Future Work

- Implement community contribution workflow
  - Evaluation
  - Validation
- GUI for DataStreamCLI
- Stay up-to-date with hydrofabric
- Academic article

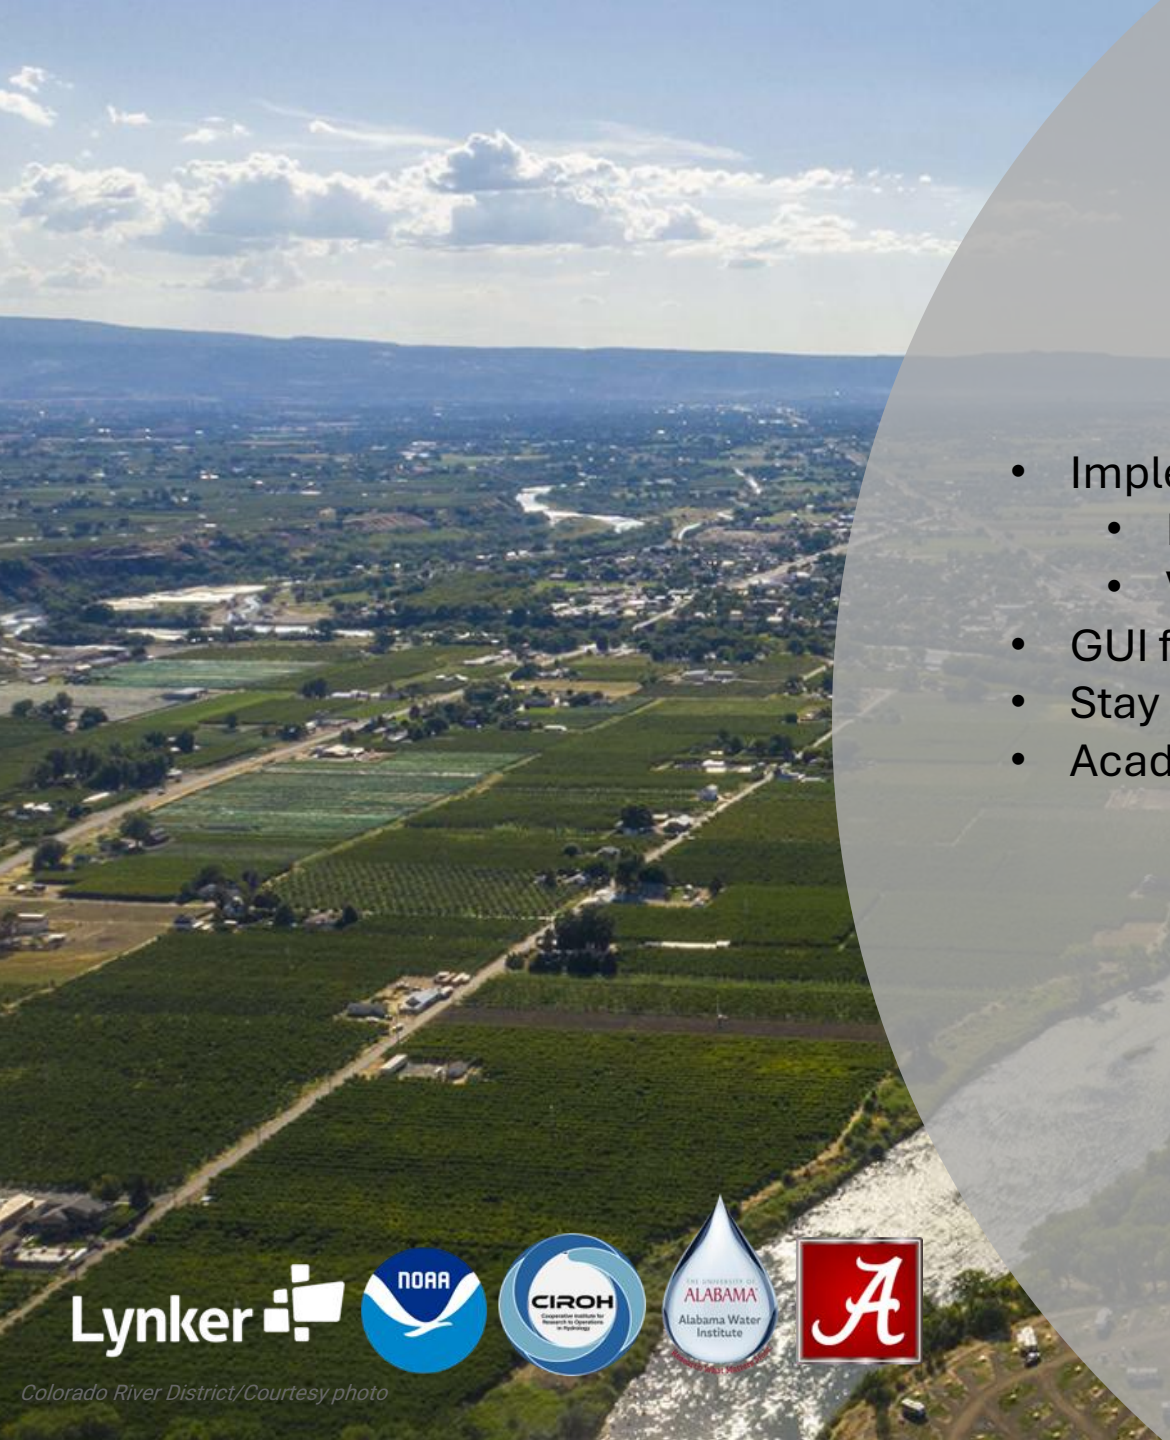

Colorado River District/Courtesy photo